

Assignment1_kwok

January 10, 2024

```
[1]: #Import libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

# Set Pandas options to show all columns and rows
pd.set_option('display.max_rows', None)
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)
pd.set_option('display.max_colwidth', None)

# Make variable for input file
INFILE = "/Users/jck/Documents/MSDS 422/Unit 1/Assignment 1/HMEQ_Loss.csv"

# Read in the data file
df = pd.read_csv(INFILE, sep=',', header=0)
```

```
[2]: # Print the first 5 rows of the data frame
df.head(5)
```

```
[2]:
```

	TARGET_BAD_FLAG	TARGET_LOSS_AMT	LOAN	MORTDUE	VALUE	REASON	JOB	\
0	1	641.0	1100	25860.0	39025.0	HomeImp	Other	
1	1	1109.0	1300	70053.0	68400.0	HomeImp	Other	
2	1	767.0	1500	13500.0	16700.0	HomeImp	Other	
3	1	1425.0	1500	NaN	NaN	NaN	NaN	
4	0	NaN	1700	97800.0	112000.0	HomeImp	Office	

	YOJ	DEROG	DELINQ	CLAGE	NINQ	CLNO	DEBTINC
0	10.5	0.0	0.0	94.366667	1.0	9.0	NaN
1	7.0	0.0	2.0	121.833333	0.0	14.0	NaN
2	4.0	0.0	0.0	149.466667	1.0	10.0	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	3.0	0.0	0.0	93.333333	0.0	14.0	NaN

```
[3]: # Print the data shape, such as how many rows and columns
print(df.shape)
```

```
(5960, 14)
```

```
[4]: # Print the data types for each column
dt = df.dtypes
print(dt)
```

```
TARGET_BAD_FLAG      int64
TARGET_LOSS_AMT      float64
LOAN                  int64
MORTDUE              float64
VALUE                float64
REASON               object
JOB                  object
YOJ                  float64
DEROG                float64
DELINQ               float64
CLAGE                float64
NINQ                 float64
CLNO                 float64
DEBTINC              float64
dtype: object
```

```
[5]: # Print the data frame statistics
print(df.describe())
```

	TARGET_BAD_FLAG	TARGET_LOSS_AMT	LOAN	MORTDUE	\
count	5960.000000	1189.000000	5960.000000	5442.000000	
mean	0.199497	13414.576955	18607.969799	73760.817200	
std	0.399656	10839.455965	11207.480417	44457.609458	
min	0.000000	224.000000	1100.000000	2063.000000	
25%	0.000000	5639.000000	11100.000000	46276.000000	
50%	0.000000	11003.000000	16300.000000	65019.000000	
75%	0.000000	17634.000000	23300.000000	91488.000000	
max	1.000000	78987.000000	89900.000000	399550.000000	

	VALUE	YOJ	DEROG	DELINQ	CLAGE	\
count	5848.000000	5445.000000	5252.000000	5380.000000	5652.000000	
mean	101776.048741	8.922268	0.254570	0.449442	179.766275	
std	57385.775334	7.573982	0.846047	1.127266	85.810092	
min	8000.000000	0.000000	0.000000	0.000000	0.000000	
25%	66075.500000	3.000000	0.000000	0.000000	115.116702	
50%	89235.500000	7.000000	0.000000	0.000000	173.466667	
75%	119824.250000	13.000000	0.000000	0.000000	231.562278	
max	855909.000000	41.000000	10.000000	15.000000	1168.233561	

	NINQ	CLNO	DEBTINC
--	------	------	---------

count	5450.000000	5738.000000	4693.000000
mean	1.186055	21.296096	33.779915
std	1.728675	10.138933	8.601746
min	0.000000	0.000000	0.524499
25%	0.000000	15.000000	29.140031
50%	1.000000	20.000000	34.818262
75%	2.000000	26.000000	39.003141
max	17.000000	71.000000	203.312149

```
[6]: # Print the number of missing values for each column
missing_values = df.isnull().sum()
print(missing_values)
# After running the above code, we can see that there are 12 columns with
↳ missing values.
```

```
TARGET_BAD_FLAG      0
TARGET_LOSS_AMT     4771
LOAN                  0
MORTDUE              518
VALUE                112
REASON               252
JOB                  279
YOJ                  515
DEROG                708
DELINQ              580
CLAGE                308
NINQ                 510
CLNO                 222
DEBTINC             1267
dtype: int64
```

```
[7]: # Show which column is under object type and which is under numeric type
TARGET_B = "TARGET_BAD_FLAG"
TARGET_L = "TARGET_LOSS_AMT"

objList = []
numList = []
for i in dt.index :
    #print(" here is i .....", i , " ..... and here is the type", dt[i] )
    if i in ( [ TARGET_B, TARGET_L ] ) : continue
    if dt[i] in (["object"]) : objList.append( i )
    if dt[i] in (["float64","int64"]) : numList.append( i )

print(" OBJECTS ")
print(" ----- ")
for i in objList :
    print( i )
print(" ----- ")
```

```

print(" NUMBER ")
print(" ----- ")
for i in numList :
    print( i )
print(" ----- ")

```

```

OBJECTS
-----
REASON
JOB
-----
NUMBER
-----
LOAN
MORTDUE
VALUE
YOJ
DEROG
DELINQ
CLAGE
NINQ
CLNO
DEBTINC
-----

```

```

[8]: # My idea is insert 0 for Target Loss Amount for my first step of data cleaning
df[TARGET_L] = df[TARGET_L].fillna(0)

```

```

missing_values2 =df.isnull().sum()
print(missing_values2)
'''

```

*After setting the Target Loss Amount to 0, there are no longer any missing
 ↪values in that column.*

*I chose not to impute values for Target Loss Amount because it should be 0 if
 ↪the loan was repaid successfully.*

Consequently, the count of columns with missing values has decreased to 11.
 '''

```

TARGET_BAD_FLAG      0
TARGET_LOSS_AMT      0
LOAN                  0
MORTDUE               518
VALUE                 112
REASON                252
JOB                   279
YOJ                   515

```

```

DEROG          708
DELINQ         580
CLAGE          308
NINQ           510
CLNO           222
DEBTINC        1267
dtype: int64

```

[8]: '\nAfter setting the Target Loss Amount to 0, there are no longer any missing values in that column. \nI chose not to impute values for Target Loss Amount because it should be 0 if the loan was repaid successfully. \nConsequently, the count of columns with missing values has decreased to 11.\n'

[9]: *# Next, lets handle the numeric columns with missing values*

```

cols_with_missing = ['MORTDUE', 'VALUE', 'YOJ', 'DEROG', 'DELINQ', 'CLAGE',
↳ 'NINQ', 'CLNO', 'DEBTINC']

```

[10]:

```

'''
I observed that some data points are outliers, so I plan to employ the
↳ Interquartile Range (IQR) method
to detect these outliers and substitute them, using the median value for
↳ imputation.
'''
for col in cols_with_missing:
    # 1. Identify Outliers using the IQR method
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # 2. Remove Outliers - Replace outliers in a copy of the column with NaN
    temp_col = df[col].copy()
    temp_col[(temp_col < lower_bound) | (temp_col > upper_bound)] = np.nan

    # 3. Calculate median of the column with outliers removed
    median_val = temp_col.median()

    # 4. Create new column for imputed values, fill missing values with the
    ↳ calculated median
    df['IMP_'+col] = df[col].fillna(median_val)

```

[11]: `print(df.head(5))`

	TARGET_BAD_FLAG	TARGET_LOSS_AMT	LOAN	MORTDUE	VALUE	REASON	JOB	\
0	1	641.0	1100	25860.0	39025.0	HomeImp	Other	
1	1	1109.0	1300	70053.0	68400.0	HomeImp	Other	

2	1	767.0	1500	13500.0	16700.0	HomeImp	Other
3	1	1425.0	1500	NaN	NaN	NaN	NaN
4	0	0.0	1700	97800.0	112000.0	HomeImp	Office

	YOJ	DEROG	DELINQ	CLAGE	NINQ	CLNO	DEBTINC	IMP_MORTDUE	\
0	10.5	0.0	0.0	94.366667	1.0	9.0	NaN	25860.0	
1	7.0	0.0	2.0	121.833333	0.0	14.0	NaN	70053.0	
2	4.0	0.0	0.0	149.466667	1.0	10.0	NaN	13500.0	
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	63508.0	
4	3.0	0.0	0.0	93.333333	0.0	14.0	NaN	97800.0	

	IMP_VALUE	IMP_YOJ	IMP_DEROG	IMP_DELINQ	IMP_CLAGE	IMP_NINQ	IMP_CLNO	\
0	39025.0	10.5	0.0	0.0	94.366667	1.0	9.0	
1	68400.0	7.0	0.0	2.0	121.833333	0.0	14.0	
2	16700.0	4.0	0.0	0.0	149.466667	1.0	10.0	
3	86908.0	7.0	0.0	0.0	172.432355	1.0	20.0	
4	112000.0	3.0	0.0	0.0	93.333333	0.0	14.0	

	IMP_DEBTINC
0	34.880462
1	34.880462
2	34.880462
3	34.880462
4	34.880462

```
[12]: missing_values3 =df.isnull().sum()
print(missing_values3)
# After running the above code, we can see that there are 2 columns with
↳missing values.
```

TARGET_BAD_FLAG	0
TARGET_LOSS_AMT	0
LOAN	0
MORTDUE	518
VALUE	112
REASON	252
JOB	279
YOJ	515
DEROG	708
DELINQ	580
CLAGE	308
NINQ	510
CLNO	222
DEBTINC	1267
IMP_MORTDUE	0
IMP_VALUE	0
IMP_YOJ	0
IMP_DEROG	0

```

IMP_DELINQ          0
IMP_CLAGE           0
IMP_NINQ            0
IMP_CLNO            0
IMP_DEBTINC         0
dtype: int64

```

```

[13]: # Let's now address the missing values in the categorical columns.
      categorical_cols_with_missing = ['REASON', 'JOB']

```

```

[14]: for col in categorical_cols_with_missing :
      # 1. Fill Missing Values as a separate category
      df['TEMP_'+col] = df[col].fillna('Missing')

      # 2. One-Hot Encode in the temporary column
      encoded = pd.get_dummies(df['TEMP_'+col], prefix='OHE_'+col, dtype=int)
      # print(encoded.head(5))
      # 3. Merge the new one-hot encoded columns back with df
      df = pd.concat([df, encoded], axis=1)

      # 4. Remove the temporary column
      df.drop(columns='TEMP_'+col, inplace=True)

      # After this loop, 'df' will have new one-hot encoded columns corresponding to
      ↪ each category in the original columns.

```

```

[15]: # Validate that all missing values have been handled
      missing_values4 = df.isnull().sum()
      print(missing_values4)

```

```

TARGET_BAD_FLAG      0
TARGET_LOSS_AMT      0
LOAN                 0
MORTDUE              518
VALUE                112
REASON               252
JOB                  279
YOJ                  515
DEROG                708
DELINQ               580
CLAGE                308
NINQ                 510
CLNO                 222
DEBTINC              1267
IMP_MORTDUE           0
IMP_VALUE             0
IMP_YOJ               0
IMP_DEROG             0

```

```

IMP_DELINQ          0
IMP_CLAGE           0
IMP_NINQ            0
IMP_CLNO            0
IMP_DEBTINC         0
OHE_REASON_DebtCon  0
OHE_REASON_HomeImp  0
OHE_REASON_Missing  0
OHE_JOB_Mgr         0
OHE_JOB_Missing     0
OHE_JOB_Office      0
OHE_JOB_Other       0
OHE_JOB_ProfExe     0
OHE_JOB_Sales       0
OHE_JOB_Self        0
dtype: int64

```

```

[16]: print(missing_values4.tail(10))
      # After running the code above, I noticed we have no more missing value

```

```

OHE_REASON_DebtCon  0
OHE_REASON_HomeImp  0
OHE_REASON_Missing  0
OHE_JOB_Mgr         0
OHE_JOB_Missing     0
OHE_JOB_Office      0
OHE_JOB_Other       0
OHE_JOB_ProfExe     0
OHE_JOB_Sales       0
OHE_JOB_Self        0
dtype: int64

```

```

[17]: df.describe()

```

```

[17]:      TARGET_BAD_FLAG  TARGET_LOSS_AMT      LOAN      MORTDUE  \
count      5960.000000      5960.000000  5960.000000  5442.000000
mean         0.199497      2676.163087  18607.969799  73760.817200
std          0.399656      7222.631500  11207.480417  44457.609458
min           0.000000         0.000000   1100.000000   2063.000000
25%           0.000000         0.000000  11100.000000  46276.000000
50%           0.000000         0.000000  16300.000000  65019.000000
75%           0.000000         0.000000  23300.000000  91488.000000
max           1.000000      78987.000000  89900.000000 399550.000000

      VALUE      YOJ      DEROG      DELINQ      CLAGE  \
count  5848.000000  5445.000000  5252.000000  5380.000000  5652.000000
mean  101776.048741    8.922268    0.254570    0.449442   179.766275
std   57385.775334    7.573982    0.846047    1.127266   85.810092

```


min	8000.000000	0.000000	0.000000	0.000000	0.000000
25%	66075.500000	3.000000	0.000000	0.000000	115.116702
50%	89235.500000	7.000000	0.000000	0.000000	173.466667
75%	119824.250000	13.000000	0.000000	0.000000	231.562278
max	855909.000000	41.000000	10.000000	15.000000	1168.233561

	NINQ	CLNO	DEBTINC	IMP_MORTDUE	IMP_VALUE \
count	5450.000000	5738.000000	4693.000000	5960.000000	5960.000000
mean	1.186055	21.296096	33.779915	72869.716644	101496.649168
std	1.728675	10.138933	8.601746	42579.485794	56879.779380
min	0.000000	0.000000	0.524499	2063.000000	8000.000000
25%	0.000000	15.000000	29.140031	48139.000000	66489.500000
50%	1.000000	20.000000	34.818262	63508.000000	88310.500000
75%	2.000000	26.000000	39.003141	88200.250000	119004.750000
max	17.000000	71.000000	203.312149	399550.000000	855909.000000

	IMP_YOJ	IMP_DEROG	IMP_DELINQ	IMP_CLAGE	IMP_NINQ \
count	5960.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	8.756166	0.224329	0.405705	179.387274	1.170134
std	7.259424	0.798458	1.079256	83.578832	1.653866
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	3.000000	0.000000	0.000000	117.371430	0.000000
50%	7.000000	0.000000	0.000000	172.432355	1.000000
75%	12.000000	0.000000	0.000000	227.143058	2.000000
max	41.000000	10.000000	15.000000	1168.233561	17.000000

	IMP_CLNO	IMP_DEBTINC	OHE_REASON_DebtCon	OHE_REASON_HomeImp \
count	5960.000000	5960.000000	5960.000000	5960.000000
mean	21.247819	34.013874	0.659060	0.298658
std	9.951308	7.645985	0.474065	0.457708
min	0.000000	0.524499	0.000000	0.000000
25%	15.000000	30.763159	0.000000	0.000000
50%	20.000000	34.880462	1.000000	0.000000
75%	26.000000	37.949892	1.000000	1.000000
max	71.000000	203.312149	1.000000	1.000000

	OHE_REASON_Missing	OHE_JOB_Mgr	OHE_JOB_Missing	OHE_JOB_Office \
count	5960.000000	5960.000000	5960.000000	5960.000000
mean	0.042282	0.128691	0.046812	0.159060
std	0.201248	0.334886	0.211254	0.365763
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

OHE_JOB_Other	OHE_JOB_ProfExe	OHE_JOB_Sales	OHE_JOB_Self
---------------	-----------------	---------------	--------------

count	5960.000000	5960.000000	5960.000000	5960.000000
mean	0.400671	0.214094	0.018289	0.032383
std	0.490076	0.410227	0.134004	0.177029
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

1 After filling all the missing values, I want to compare the IMP value and the original value in histogram.

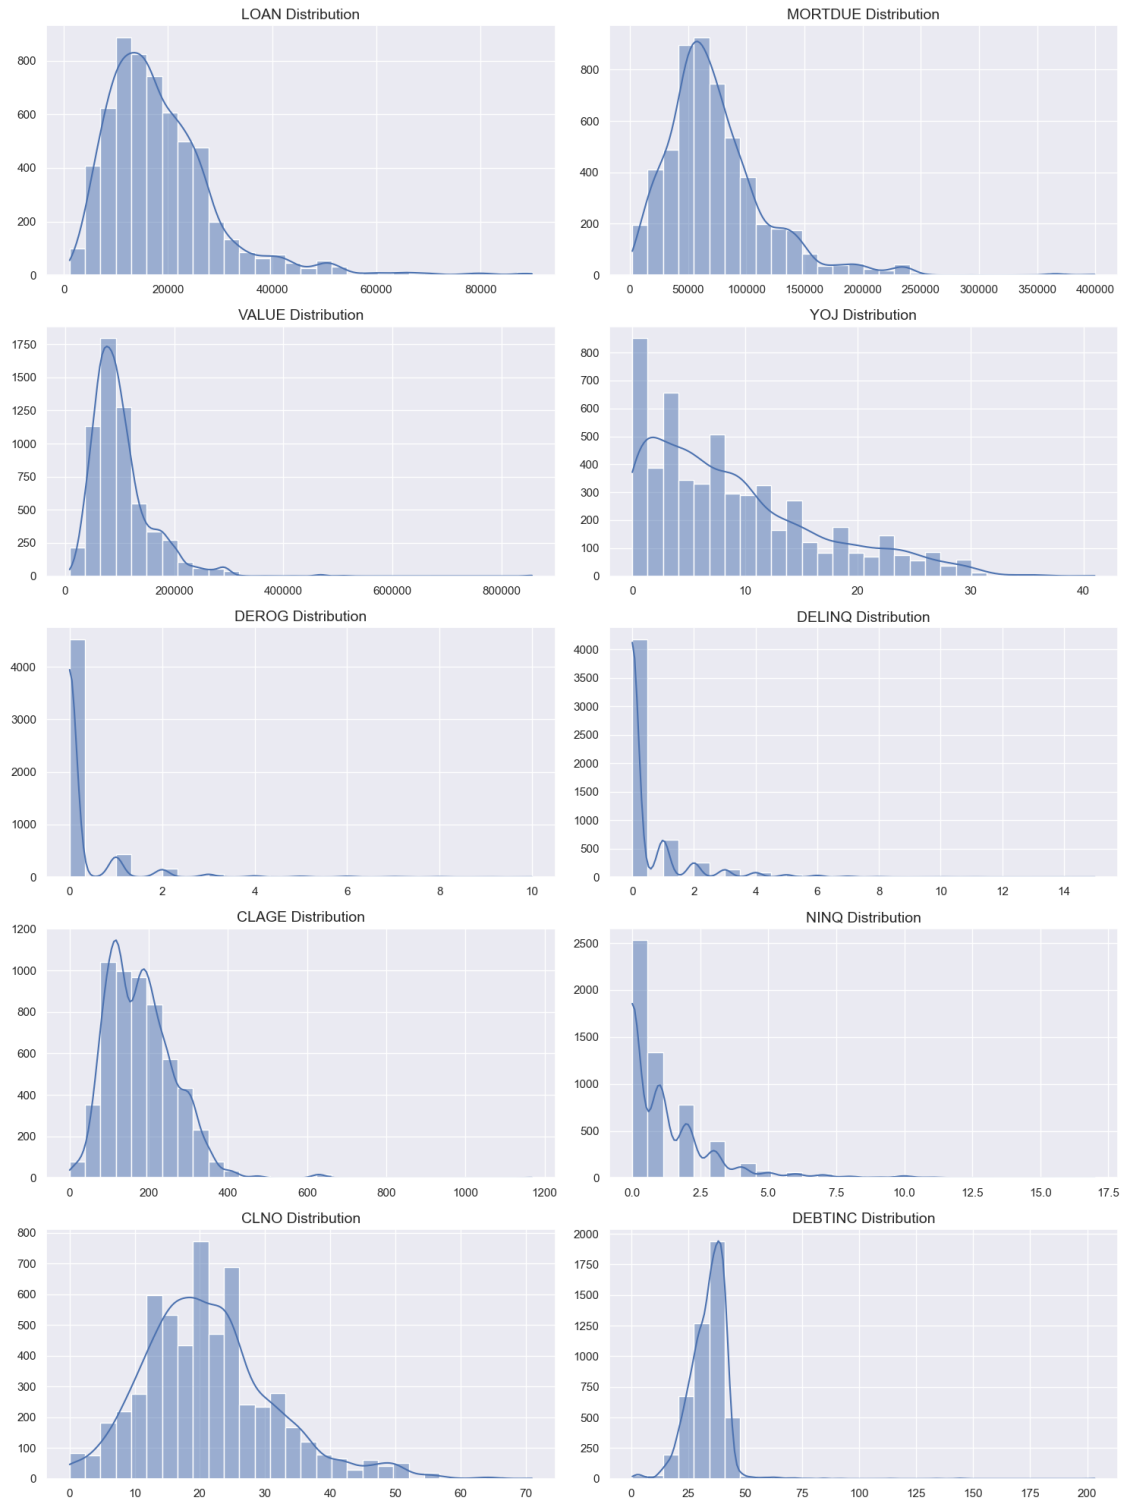
```
[18]: # The list of key numerical variables for visualization
numerical_vars = ['LOAN', 'MORTDUE', 'VALUE', 'YOJ', 'DEROG', 'DELINQ', 'CLAGE', 'NINQ', 'CLNO', 'DEBTINC']

# Initialize the subplot function using matplotlib
fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(15, 20))

# Flatten the axes array for easy iteration
axes = axes.flatten()

# Create a histogram for each numerical variable
for i, var in enumerate(numerical_vars):
    sns.histplot(df[var], bins=30, kde=True, ax=axes[i])
    axes[i].set_title(var + ' Distribution', fontsize=14)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Adjust the layout
plt.tight_layout()
plt.show()
```



```
[19]: # Define the list of key numerical variables for visualization
```

```

numerical_vars = ['LOAN', 'IMP_MORTDUE', 'IMP_VALUE', 'IMP_YOJ', 'IMP_DEROG', 'IMP_DEBTINC',
                  'IMP_DELINQ', 'IMP_CLAGE', 'IMP_NINQ', 'IMP_CLNO', 'IMP_DEBTINC']

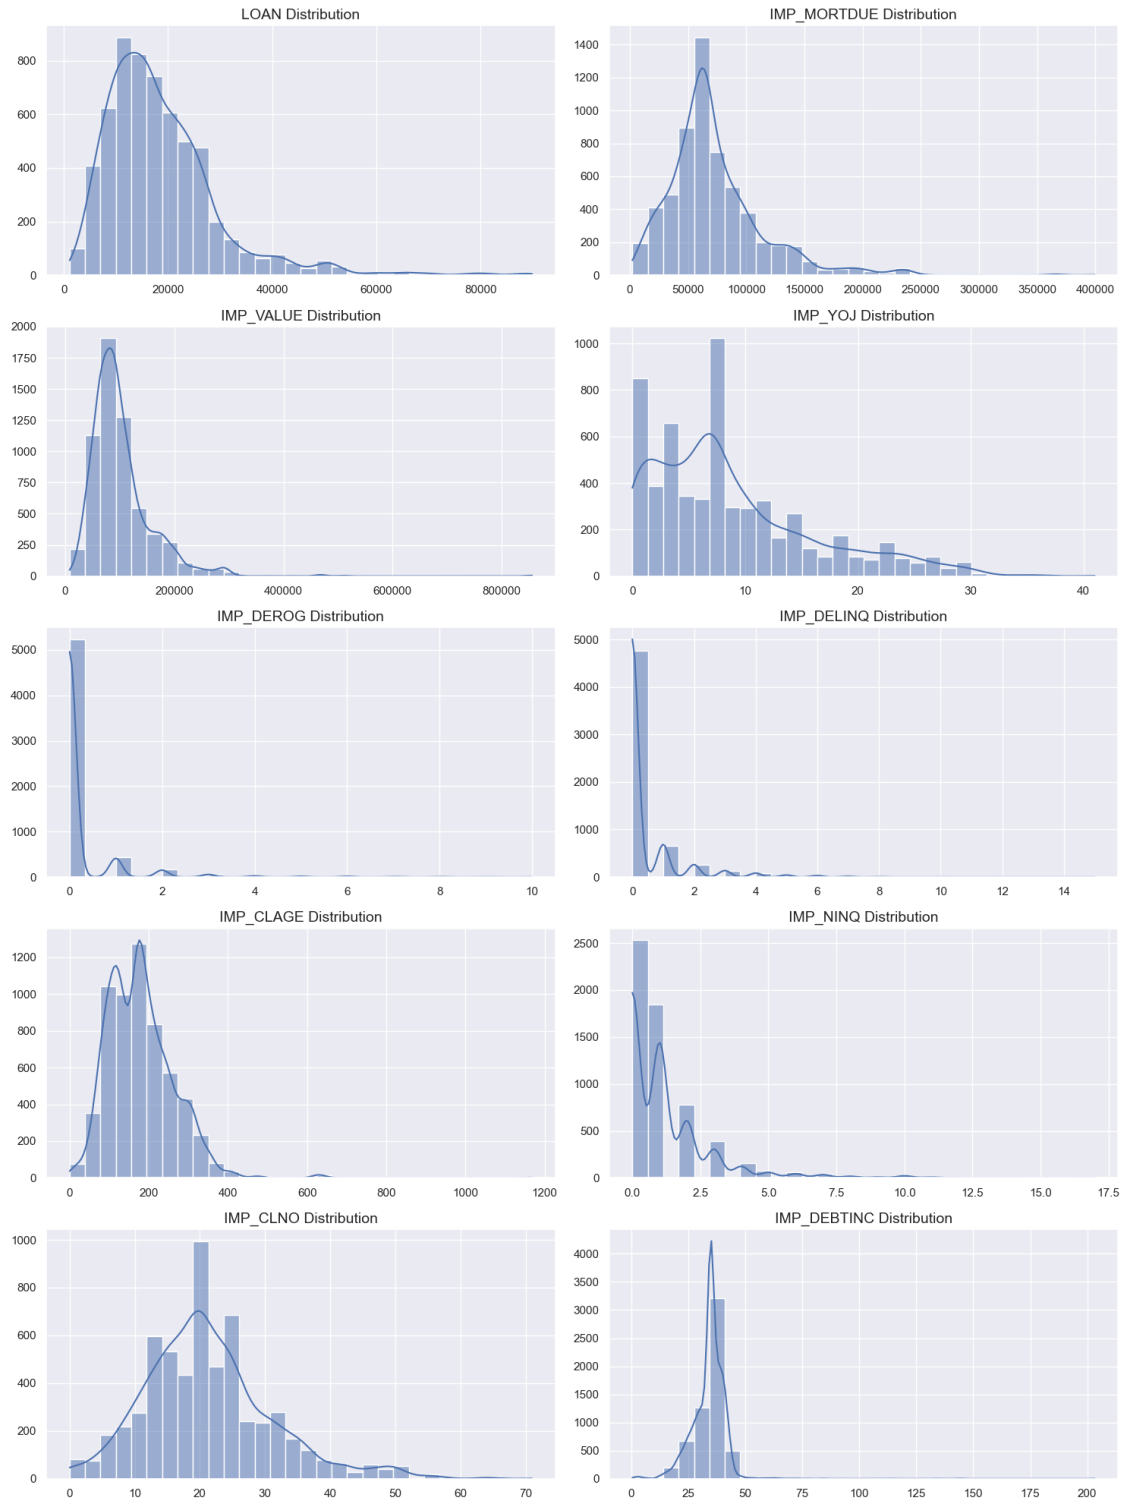
# Initialize the subplot function using matplotlib
fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(15, 20))

# Flatten the axes array for easy iteration
axes = axes.flatten()

# Create a histogram for each numerical variable
for i, var in enumerate(numerical_vars):
    sns.histplot(df[var], bins=30, kde=True, ax=axes[i])
    axes[i].set_title(var + ' Distribution', fontsize=14)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Adjust the layout
plt.tight_layout()
plt.show()

```



[20]: '''

Here are the observed findings from the histograms after filling missing values:

```

All the variables, including Loan Amount (LOAN), Mortgage Due (MORTDUE),
↳Property Value (VALUE),
Years of Employment (YOJ), Derogatory Reports (DEROG), Delinquent Credit Lines
↳(DELINQ),
Age of Oldest Credit Line (CLAGE), Number of Recent Credit Lines (NINQ), Number
↳of Credit Lines (CLNO)
, and Debt-to-Income Ratio (DEBTINC) show right-skewed distributions.
This means they mostly bunch up on the lower end with fewer instances
↳stretching out towards higher values.
The process of filling in missing data appears to have preserved the overall
↳shape of these distributions,
with no significant changes in their values.
'''

```

```

[20]: '\nHere are the observed findings from the histograms after filling missing
values:\n\nAll the variables, including Loan Amount (LOAN), Mortgage Due
(MORTDUE), Property Value (VALUE), \nYears of Employment (YOJ), Derogatory
Reports (DEROG), Delinquent Credit Lines (DELINQ), \nAge of Oldest Credit Line
(CLAGE), Number of Recent Credit Lines (NINQ), Number of Credit Lines (CLNO)\n,
and Debt-to-Income Ratio (DEBTINC) show right-skewed distributions. \nThis means
they mostly bunch up on the lower end with fewer instances stretching out
towards higher values. \nThe process of filling in missing data appears to have
preserved the overall shape of these distributions, \nwith no significant
changes in their values.\n'

```

```

[21]: '''
I want to create a new dataframe with just the filled-in columns,
and I'll make sure it doesn't have any columns with missing values , object
↳types, and target values.
'''

TARGET_COLUMNS = ['TARGET_BAD_FLAG', 'TARGET_LOSS_AMT']

# Identify numeric columns (exclude object type)
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()

# Exclude target columns from numeric columns
non_target_numeric_cols = [col for col in numeric_cols if col not in
↳TARGET_COLUMNS]

# Exclude columns with any missing values
final_cols = [col for col in non_target_numeric_cols if not df[col].isnull().
↳any()]

# Create a new DataFrame df2 with the specified columns from df
df2 = df[final_cols].copy()

```

```
# Now df2 contains only non-target, non-object, and non-missing value columns
↳ from df
```

```
df2.head(5)
```

```
[21]:
```

	LOAN	IMP_MORTDUE	IMP_VALUE	IMP_YOJ	IMP_DEROG	IMP_DELIQ	IMP_CLAGE	\
0	1100	25860.0	39025.0	10.5	0.0	0.0	94.366667	
1	1300	70053.0	68400.0	7.0	0.0	2.0	121.833333	
2	1500	13500.0	16700.0	4.0	0.0	0.0	149.466667	
3	1500	63508.0	86908.0	7.0	0.0	0.0	172.432355	
4	1700	97800.0	112000.0	3.0	0.0	0.0	93.333333	

	IMP_NINQ	IMP_CLNO	IMP_DEBTINC	OHE_REASON_DebtCon	OHE_REASON_HomeImp	\
0	1.0	9.0	34.880462	0	1	
1	0.0	14.0	34.880462	0	1	
2	1.0	10.0	34.880462	0	1	
3	1.0	20.0	34.880462	0	0	
4	0.0	14.0	34.880462	0	1	

	OHE_REASON_Missing	OHE_JOB_Mgr	OHE_JOB_Missing	OHE_JOB_Office	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	1	0	1	0	
4	0	0	0	1	

	OHE_JOB_Other	OHE_JOB_ProfExe	OHE_JOB_Sales	OHE_JOB_Self
0	1	0	0	0
1	1	0	0	0
2	1	0	0	0
3	0	0	0	0
4	0	0	0	0

```
[22]: '''
After creating the df2 dataframe,
I want to calculate the correlation between each variable and the target
↳ variable 'TARGET_LOSS_AMT'
'''

# Calculate the correlation of df2's variables with TARGET_LOSS_AMT
correlation_with_target_loss = df2.corrwith(df['TARGET_LOSS_AMT']).
↳ sort_values(ascending=False)

# Calculate the correlation of each column in df2 with TARGET_LOSS_AMT from df
correlations = df2.apply(lambda col: col.corr(df['TARGET_LOSS_AMT']))
sorted_correlations = correlations.sort_values(ascending=False)
```

```

# Print each column's correlation with TARGET_LOSS_AMT
print("Correlation of each column in df2 with TARGET_LOSS_AMT:")
print(sorted_correlations)

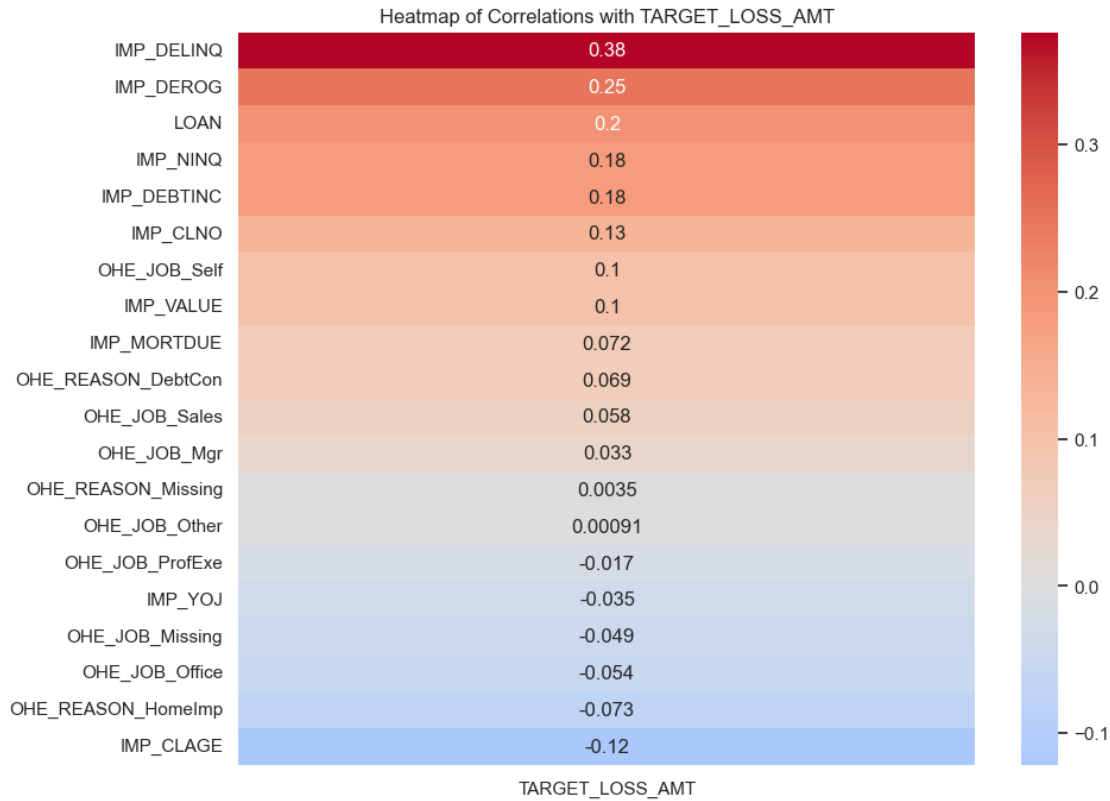
# I am trying to create a heatmap of the correlations with TARGET_LOSS_AMT
correlation_matrix = correlation_with_target_loss.
    ↳to_frame(name='TARGET_LOSS_AMT')
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Heatmap of Correlations with TARGET_LOSS_AMT')
plt.show()

```

Correlation of each column in df2 with TARGET_LOSS_AMT:

IMP_DELINQ	0.376198
IMP_DEROG	0.249934
LOAN	0.199414
IMP_NINQ	0.183584
IMP_DEBTINC	0.182208
IMP_CLNO	0.134767
OHE_JOB_Self	0.101451
IMP_VALUE	0.100304
IMP_MORTDUE	0.071745
OHE_REASON_DebtCon	0.069198
OHE_JOB_Sales	0.057618
OHE_JOB_Mgr	0.033213
OHE_REASON_Missing	0.003464
OHE_JOB_Other	0.000908
OHE_JOB_ProfExe	-0.017369
IMP_YOJ	-0.034983
OHE_JOB_Missing	-0.048824
OHE_JOB_Office	-0.054159
OHE_REASON_HomeImp	-0.073194
IMP_CLAGE	-0.121376

dtype: float64



[23]: '''

The heatmap provided displays the correlation coefficients between various imputed variables and the TARGET_LOSS_AMT. Here are the findings:

Positive Correlations: Most variables show a positive correlation with TARGET_LOSS_AMT, meaning as their values increase, the loss amount tends to increase as well. The variable IMP_DELINQ has the strongest positive correlation at 0.38, followed by IMP_DEROG at 0.25, suggesting that an increase in delinquent credit lines or derogatory reports is associated with a higher loss amount.

Negative Correlations: Two variables, IMP_YOJ and IMP_CLAGE, show a negative correlation with TARGET_LOSS_AMT (at -0.035 and -0.12 respectively), indicating that as the years of employment and the age of the oldest credit line increase, the loss amount tends to decrease.

Weak Correlations: Some variables like IMP_MORTDUE have a very low positive correlation (0.072)

'''

[23]: '\n\nThe heatmap provided displays the correlation coefficients between various imputed variables and the TARGET_LOSS_AMT. Here are the findings:\n\nPositive Correlations: Most variables show a positive correlation with TARGET_LOSS_AMT, meaning as their values increase, the loss amount tends to increase as well. The variable IMP_DELIHQ has the strongest positive correlation at 0.38, followed by IMP_DEROG at 0.25, suggesting that an increase in delinquent credit lines or derogatory reports is associated with a higher loss amount.\n\nNegative Correlations: Two variables, IMP_YOJ and IMP_CLAGE, show a negative correlation with TARGET_LOSS_AMT (at -0.035 and -0.12 respectively), indicating that as the years of employment and the age of the oldest credit line increase, the loss amount tends to decrease.\n\nWeak Correlations: Some variables like IMP_MORTDUE have a very low positive correlation (0.072) \n'

2 BINGO BONUS WORK

[24]: '''
I am trying to fill in the missing value with different methods and see which
one is better

[Column Name]+ _Missing is replacing with Median
[Column Name]+ _Missing2 is replacing with Mean
IMP_+[Column Name] is remove outliers and replace with Median
IMP2_+[Column Name] is remove outliers and replace with Mean
'''

[24]: '\n\nI am trying to fill in the missing value with different methods and see which one is better\n\n[Column Name]+ _Missing is replacing with Median \n[Column Name]+ _Missing2 is replacing with Mean\nIMP_+[Column Name] is remove outliers and replace with Median\nIMP2_+[Column Name] is remove outliers and replace with Mean\n\n'

[25]: #[Column Name]+ _Missing is replacing with Median
Create flag columns and impute missing values with median
for col in cols_with_missing:
df[col+'_MISSING'] = df[col].fillna(df[col].median(), inplace=False)

[26]: #[Column Name]+ _Missing2 is replacing with Mean
Create flag columns and impute missing values with mean
for col in cols_with_missing:
df[col+'_MISSING2'] = df[col].fillna(df[col].mean(), inplace=False)

[27]: #IMP2_+[Column Name] is remove outliers and replace with Mean

for col in cols_with_missing:
1. Identify Outliers using the IQR method
Q1 = df[col].quantile(0.25)

```

Q3 = df[col].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# 2. Remove Outliers - Replace outliers in a copy of the column with NaN
temp_col = df[col].copy()
temp_col[(temp_col < lower_bound) | (temp_col > upper_bound)] = np.nan

# 3. Calculate mean of the column with outliers removed
median_val = temp_col.mean()

# 4. Create new column for imputed values, fill missing values with the
↳calculated median
df['IMP2_'+col] = df[col].fillna(median_val)

```

```

[28]: # One of the columns with missing values that I want to analyze
cols_with_missing_M = ['MORTDUE']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns
for col in cols_with_missing_M:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")

```

Descriptive statistics for MORTDUE:

	MORTDUE	MORTDUE_MISSING	MORTDUE_MISSING2	IMP_MORTDUE	\
count	5442.000000	5960.000000	5960.000000	5960.000000	
mean	73760.817200	73001.041812	73760.817200	72869.716644	
std	44457.609458	42552.726779	42481.395689	42579.485794	
min	2063.000000	2063.000000	2063.000000	2063.000000	
25%	46276.000000	48139.000000	48139.000000	48139.000000	
50%	65019.000000	65019.000000	69529.000000	63508.000000	
75%	91488.000000	88200.250000	88200.250000	88200.250000	

max	399550.000000	399550.000000	399550.000000	399550.000000
-----	---------------	---------------	---------------	---------------

	IMP2_MORTDUE
count	5960.000000
mean	73227.437618
std	42516.565166
min	2063.000000
25%	48139.000000
50%	67623.862942
75%	88200.250000
max	399550.000000

```
[29]: # One of the columns with missing values that I want to analyze
cols_with_missing_v = ['VALUE']

# Loop through each column to get and print the descriptive statistics for the
↳ original, flag, and imputed columns
for col in cols_with_missing_v:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳ imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")
```

Descriptive statistics for VALUE:

	VALUE	VALUE_MISSING	VALUE_MISSING2	IMP_VALUE \
count	5848.000000	5960.000000	5960.000000	5960.000000
mean	101776.048741	101540.387423	101776.048741	101496.649168
std	57385.775334	56869.436682	56843.931566	56879.779380
min	8000.000000	8000.000000	8000.000000	8000.000000
25%	66075.500000	66489.500000	66489.500000	66489.500000
50%	89235.500000	89235.500000	90000.000000	88310.500000
75%	119824.250000	119004.750000	119004.750000	119004.750000
max	855909.000000	855909.000000	855909.000000	855909.000000

	IMP2_VALUE
count	5960.000000

```

mean    101604.342443
std      56857.473131
min       8000.000000
25%      66489.500000
50%      90000.000000
75%     119004.750000
max     855909.000000

```

```

[30]: # One of the columns with missing values that I want to analyze
cols_with_missing_y = ['YOJ']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns
for col in cols_with_missing_y:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")

```

Descriptive statistics for YOJ:

	YOJ	YOJ_MISSING	YOJ_MISSING2	IMP_YOJ	IMP2_YOJ
count	5445.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	8.922268	8.756166	8.922268	8.756166	8.889934
std	7.573982	7.259424	7.239301	7.259424	7.240065
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	3.000000	3.000000	3.000000	3.000000	3.000000
50%	7.000000	7.000000	8.000000	7.000000	8.000000
75%	13.000000	12.000000	12.000000	12.000000	12.000000
max	41.000000	41.000000	41.000000	41.000000	41.000000

```

[31]: # One of the columns with missing values that I want to analyze
cols_with_missing_d = ['DEROG']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns

```

```

for col in cols_with_missing_d:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")

```

Descriptive statistics for DEROG:

	DEROG	DEROG_MISSING	DEROG_MISSING2	IMP_DEROG	IMP2_DEROG
count	5252.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	0.254570	0.224329	0.254570	0.224329	0.224329
std	0.846047	0.798458	0.794198	0.798458	0.798458
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	10.000000	10.000000	10.000000	10.000000	10.000000

```

[32]: # One of the columns with missing values that I want to analyze
cols_with_missing_D = ['DELINQ']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns
for col in cols_with_missing_D:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

    # Calculate and print descriptive statistics

```

```

descriptive_stats = df[columns_to_compare].describe()
print(f"Descriptive statistics for {col}:")
print(descriptive_stats, "\n")

```

Descriptive statistics for DELINQ:

	DELINQ	DELINQ_MISSING	DELINQ_MISSING2	IMP_DELINQ	IMP2_DELINQ
count	5380.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	0.449442	0.405705	0.449442	0.405705	0.405705
std	1.127266	1.079256	1.071002	1.079256	1.079256
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.449442	0.000000	0.000000
max	15.000000	15.000000	15.000000	15.000000	15.000000

```

[33]: # One of the columns with missing values that I want to analyze
cols_with_missing_c = ['CLAGE']

# Loop through each column to get and print the descriptive statistics for the
↳ original, flag, and imputed columns
for col in cols_with_missing_c:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳ imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")

```

Descriptive statistics for CLAGE:

	CLAGE	CLAGE_MISSING	CLAGE_MISSING2	IMP_CLAGE	IMP2_CLAGE
count	5652.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	179.766275	179.440725	179.766275	179.387274	179.609230
std	85.810092	83.574697	83.563059	83.578832	83.565767
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	115.116702	117.371430	117.371430	117.371430	117.371430
50%	173.466667	173.466667	178.076005	172.432355	176.727344
75%	231.562278	227.143058	227.143058	227.143058	227.143058
max	1168.233561	1168.233561	1168.233561	1168.233561	1168.233561

```
[34]: # One of the columns with missing values that I want to analyze
cols_with_missing_n = ['NINQ']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns
for col in cols_with_missing_n:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")
```

Descriptive statistics for NINQ:

	NINQ	NINQ_MISSING	NINQ_MISSING2	IMP_NINQ	IMP2_NINQ
count	5450.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	1.186055	1.170134	1.186055	1.170134	1.166905
std	1.728675	1.653866	1.653046	1.653866	1.654232
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	1.000000	1.000000	1.000000	0.962261
75%	2.000000	2.000000	2.000000	2.000000	2.000000
max	17.000000	17.000000	17.000000	17.000000	17.000000

```
[35]: # One of the columns with missing values that I want to analyze
cols_with_missing_C = ['CLNO']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns
for col in cols_with_missing_C:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col
```



```

# Select the columns to compare
columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

# Calculate and print descriptive statistics
descriptive_stats = df[columns_to_compare].describe()
print(f"Descriptive statistics for {col}:")
print(descriptive_stats, "\n")

```

Descriptive statistics for CLNO:

	CLNO	CLNO_MISSING	CLNO_MISSING2	IMP_CLNO	IMP2_CLNO
count	5738.000000	5960.000000	5960.000000	5960.000000	5960.000000
mean	21.296096	21.247819	21.296096	21.247819	21.254561
std	10.138933	9.951308	9.948280	9.951308	9.950521
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	15.000000	15.000000	15.000000	15.000000	15.000000
50%	20.000000	20.000000	21.000000	20.000000	20.181011
75%	26.000000	26.000000	26.000000	26.000000	26.000000
max	71.000000	71.000000	71.000000	71.000000	71.000000

```

[36]: # One of the columns with missing values that I want to analyze
cols_with_missing_De = ['DEBTINC']

# Loop through each column to get and print the descriptive statistics for the
↳original, flag, and imputed columns
for col in cols_with_missing_De:
    # Define the column names for original, missing flags and imputed columns
    original_col = col
    missing_flag = col + '_MISSING'
    missing2_flag = col + '_MISSING2'
    imp_col = 'IMP_' + col
    imp2_col = 'IMP2_' + col

    # Select the columns to compare
    columns_to_compare = [original_col, missing_flag, missing2_flag, imp_col,
↳imp2_col]

    # Calculate and print descriptive statistics
    descriptive_stats = df[columns_to_compare].describe()
    print(f"Descriptive statistics for {col}:")
    print(descriptive_stats, "\n")

```

Descriptive statistics for DEBTINC:

	DEBTINC	DEBTINC_MISSING	DEBTINC_MISSING2	IMP_DEBTINC	\
count	4693.000000	5960.000000	5960.000000	5960.000000	
mean	33.779915	34.000651	33.779915	34.013874	

std	8.601746	7.644528	7.632713	7.645985
min	0.524499	0.524499	0.524499	0.524499
25%	29.140031	30.763159	30.763159	30.763159
50%	34.818262	34.818262	33.779915	34.880462
75%	39.003141	37.949892	37.949892	37.949892
max	203.312149	203.312149	203.312149	203.312149

	IMP2_DEBTINC
count	5960.000000
mean	33.779285
std	7.632713
min	0.524499
25%	30.763159
50%	33.776951
75%	37.949892
max	203.312149

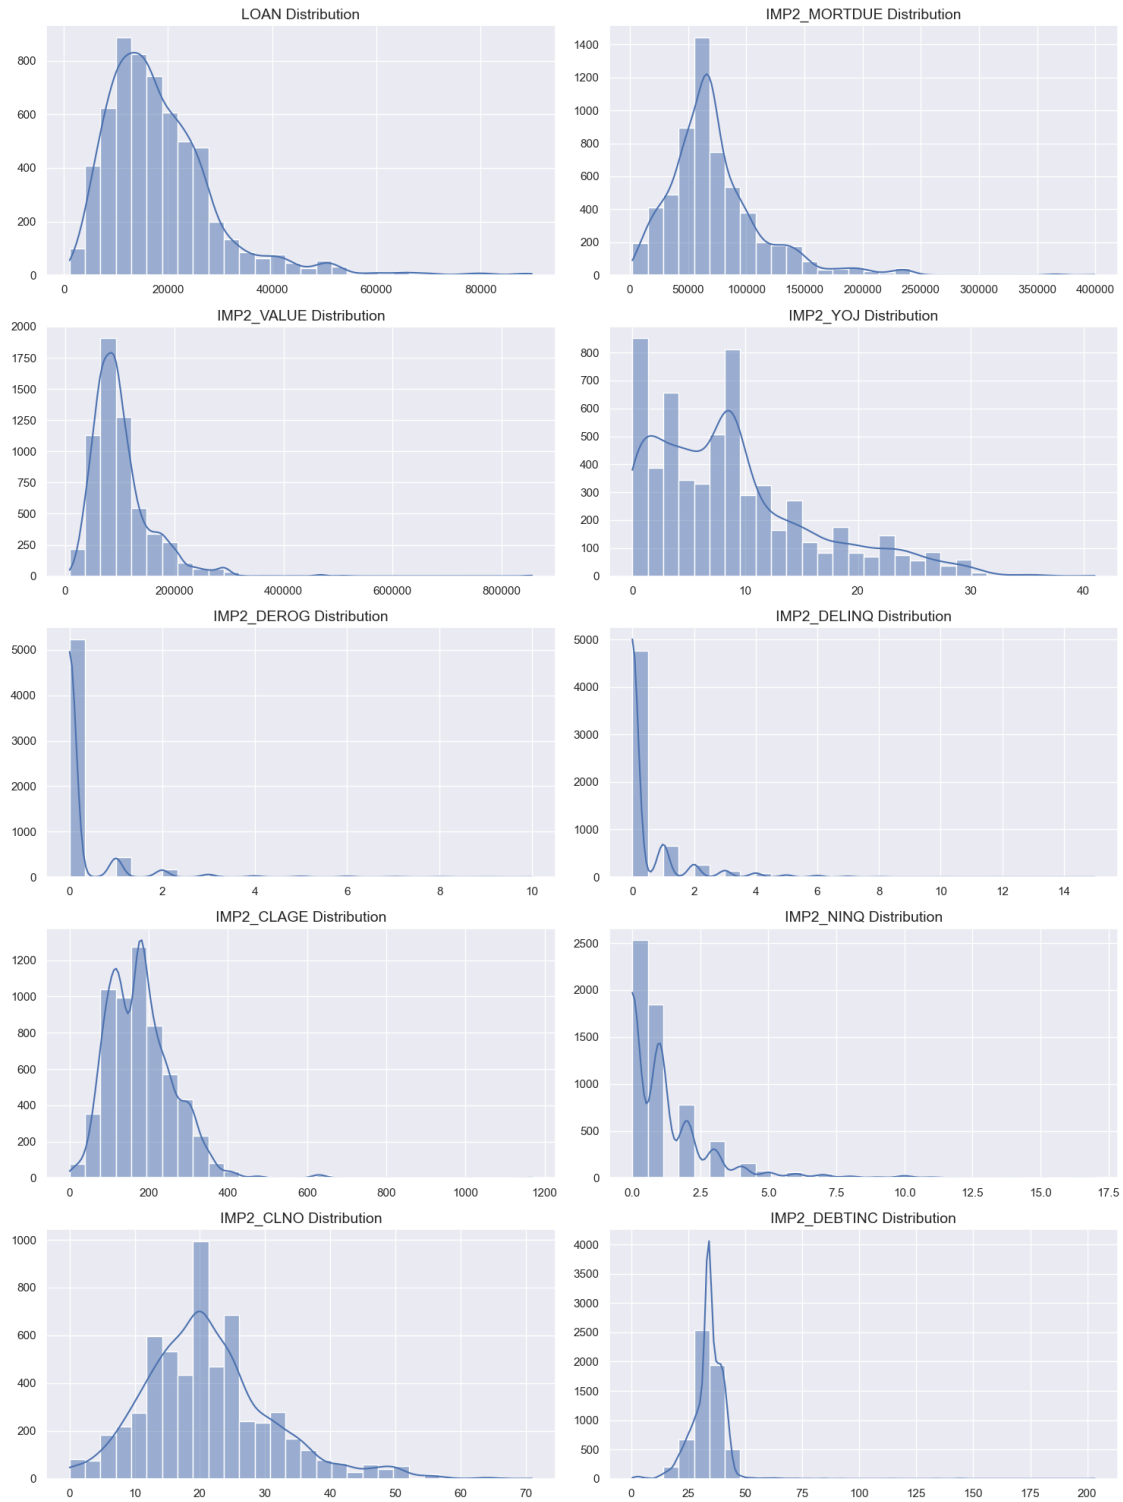
```
[37]: # Define the list of key numerical variables for visualization
numerical_vars = ['LOAN', 'IMP2_MORTDUE', 'IMP2_VALUE', 'IMP2_YOJ',
                  ↪ 'IMP2_DEROG', 'IMP2_DELIHQ', 'IMP2_CLAGE', 'IMP2_NINQ', 'IMP2_CLNO',
                  ↪ 'IMP2_DEBTINC']

# Initialize the subplot function using matplotlib
fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(15, 20))

# Flatten the axes array for easy iteration
axes = axes.flatten()

# Create a histogram for each numerical variable
for i, var in enumerate(numerical_vars):
    sns.histplot(df[var], bins=30, kde=True, ax=axes[i])
    axes[i].set_title(var + ' Distribution', fontsize=14)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Adjust the layout
plt.tight_layout()
plt.show()
```



```
[38]: # Define the list of key numerical variables for visualization
```

```

numerical_vars = ['LOAN', 'MORTDUE_MISSING', 'VALUE_MISSING', 'YOJ_MISSING',
↳ 'DEROG_MISSING', 'DELINQ_MISSING', 'CLAGE_MISSING', 'NINQ_MISSING',
↳ 'CLNO_MISSING', 'DEBTINC_MISSING']

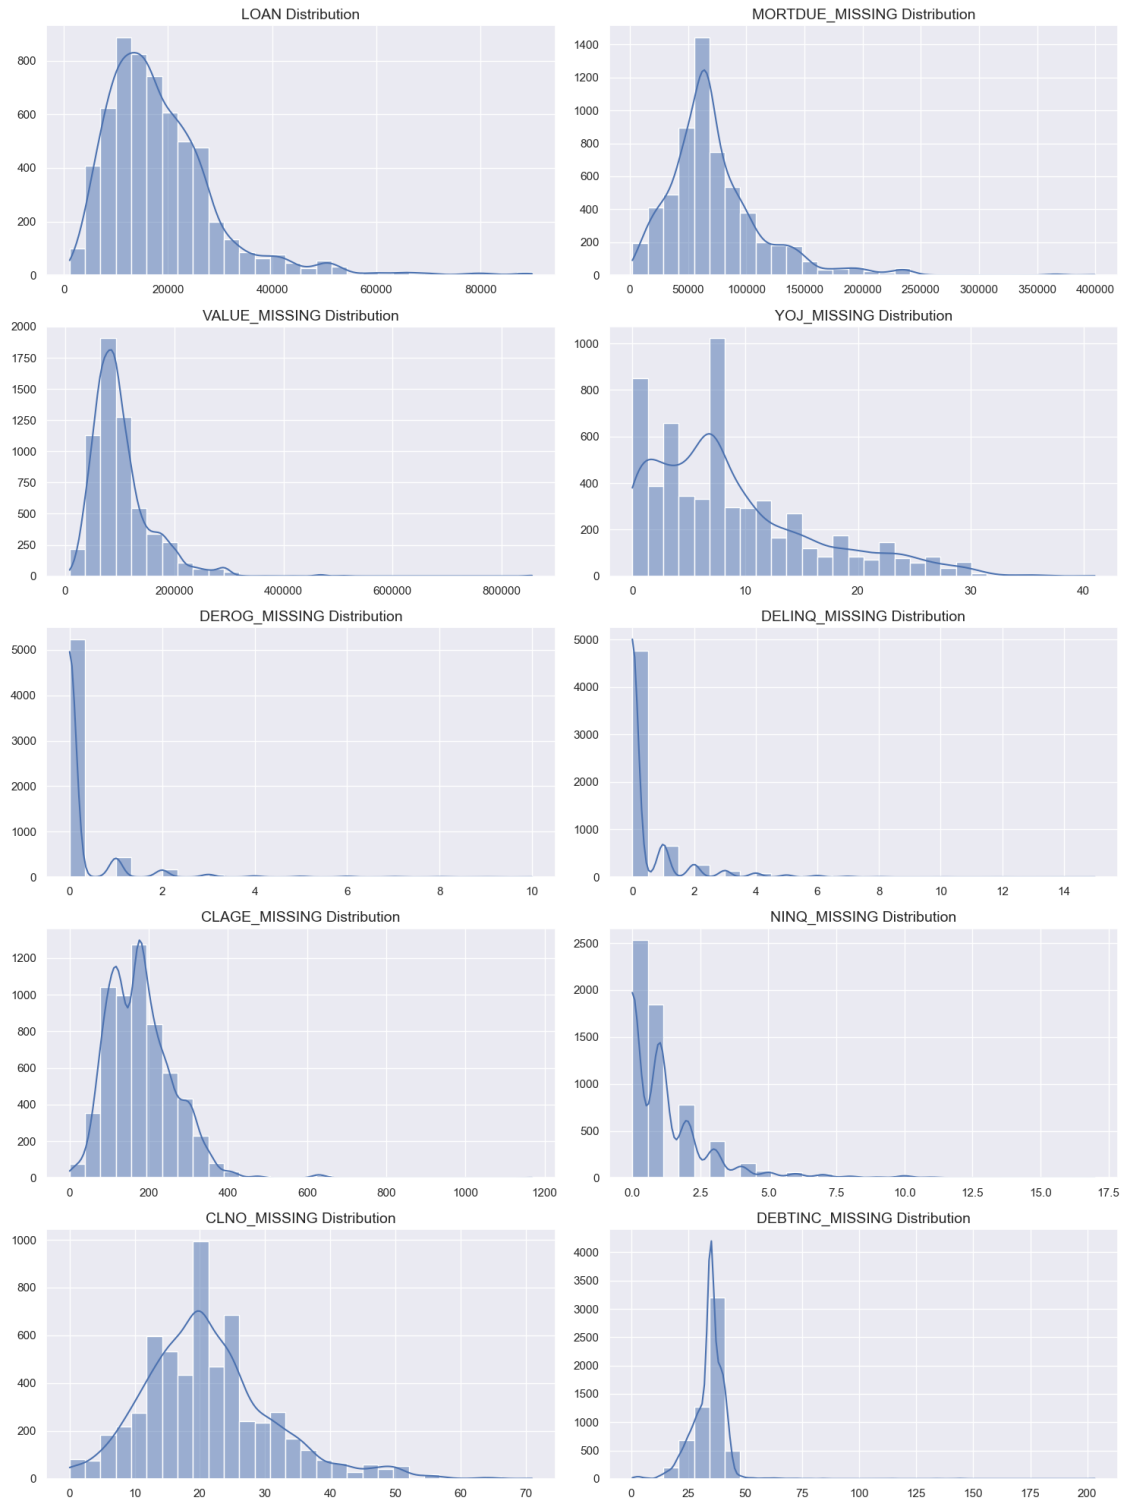
# Initialize the subplot function using matplotlib
fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(15, 20))

# Flatten the axes array for easy iteration
axes = axes.flatten()

# Create a histogram for each numerical variable
for i, var in enumerate(numerical_vars):
    sns.histplot(df[var], bins=30, kde=True, ax=axes[i])
    axes[i].set_title(var + ' Distribution', fontsize=14)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Adjust the layout
plt.tight_layout()
plt.show()

```



```
[39]: # Define the list of key numerical variables for visualization
```

```

numerical_vars = ['LOAN', 'MORTDUE_MISSING2', 'VALUE_MISSING2', 'YOJ_MISSING2',
↳ 'DEROG_MISSING2', 'DELINQ_MISSING2', 'CLAGE_MISSING2', 'NINQ_MISSING2',
↳ 'CLNO_MISSING2', 'DEBTINC_MISSING2']

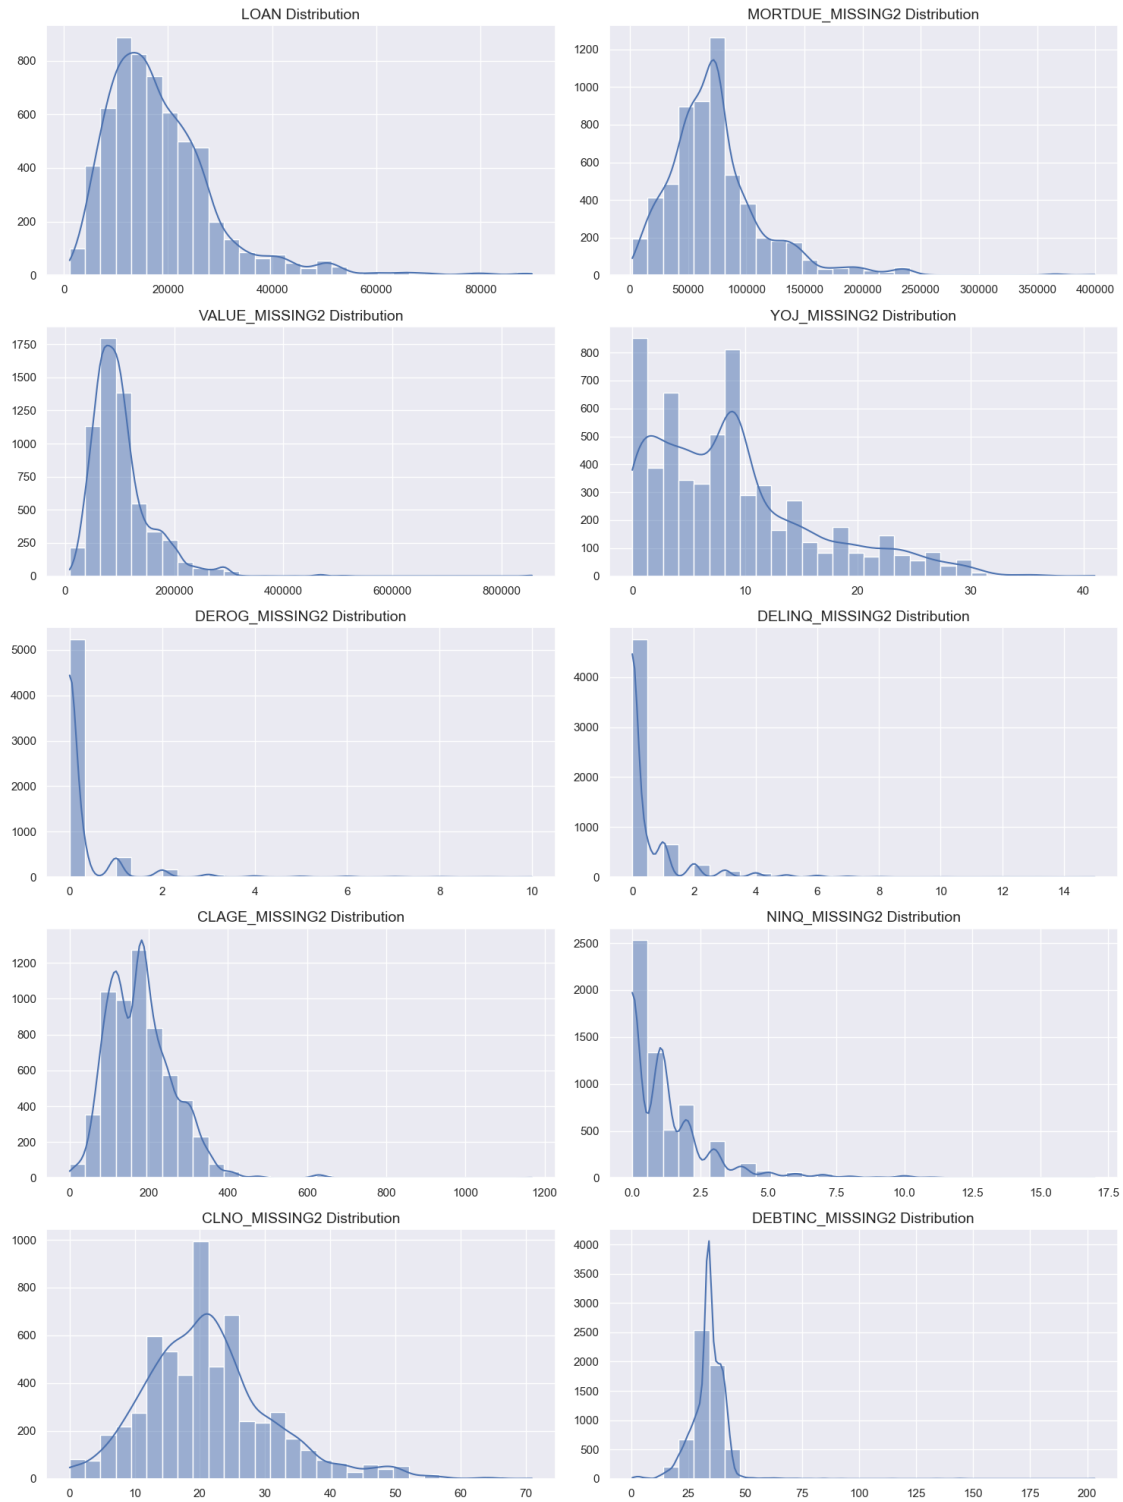
# Initialize the subplot function using matplotlib
fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(15, 20))

# Flatten the axes array for easy iteration
axes = axes.flatten()

# Create a histogram for each numerical variable
for i, var in enumerate(numerical_vars):
    sns.histplot(df[var], bins=30, kde=True, ax=axes[i])
    axes[i].set_title(var + ' Distribution', fontsize=14)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('')

# Adjust the layout
plt.tight_layout()
plt.show()

```



[40]: '''

After trying various ways to handle missing data, like:

```

[Column Name]+ _Missing is replacing with Median
[Column Name]+ _Missing2 is replacing with Mean
IMP_+[Column Name] is remove outliers and replace with Median
IMP2_+[Column Name] is remove outliers and replace with Mean
I noticed that the distributions and mean values stayed pretty similar across
↳these methods.
However, the approach where I took out the outliers and used the median,
labeled as "IMP_[Column Name]", seems to make the most sense to me.

***I've discussed with TA Logan, and he mentioned that it was acceptable,
↳despite the PDF cut off some of the output. ***
'''

```

```

[40]: '\nAfter trying various ways to handle missing data, like:\n\n[Column Name]+
_Missing is replacing with Median \n[Column Name]+ _Missing2 is replacing with
Mean\nIMP_+[Column Name] is remove outliers and replace with Median\nIMP2_
+[Column Name] is remove outliers and replace with Mean\nI noticed that the
distributions and mean values stayed pretty similar across these methods.
\nHowever, the approach where I took out the outliers and used the median,
\nlabeled as "IMP_[Column Name]", seems to make the most sense to
me.\n\n***I\'ve discussed with TA Logan, and he mentioned that it was
acceptable, despite the PDF cut off some of the output. ***\n'

```

```

[41]: !jupyter nbconvert --to pdf Assignment1_kwok.ipynb

```